

```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.discriminant_analysis import StandardScaler
from sklearn.metrics import confusion_matrix, accuracy_score,
precision_score, recall_score, f1_score, ConfusionMatrixDisplay

attributes = ['Age', 'Sector', 'idk', 'Education', 'idk2',
'Marital_Stat', 'Job', 'Family', 'Race', 'Gender', 'idk3', 'idk4',
'Hours/wk', 'Country', 'Income_>50k']

df = pd.read_csv(r"Penghasilan_Orang_Dewasa.csv", header=None)
df.columns = attributes
df.head()

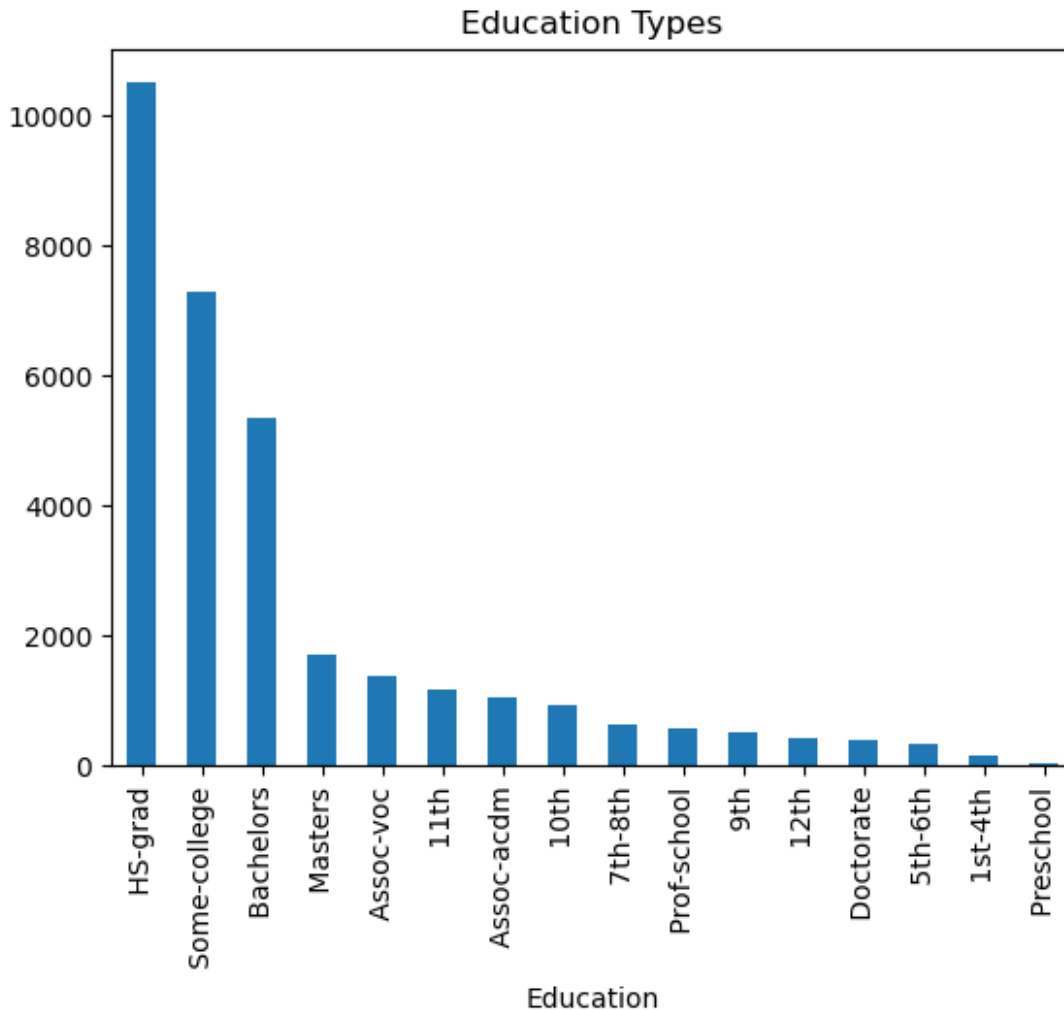
```

	Age	Sector	idk	Education	idk2	
Marital_Stat \						
0	39	State-gov	77516	Bachelors	13	Never-married
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse
2	38	Private	215646	HS-grad	9	Divorced
3	53	Private	234721	11th	7	Married-civ-spouse
4	28	Private	338409	Bachelors	13	Married-civ-spouse

	Job	Family	Race	Gender	idk3	idk4
Hours/wk \						
0	Adm-clerical	Not-in-family	White	Male	2174	0
40						
1	Exec-managerial	Husband	White	Male	0	0
13						
2	Handlers-cleaners	Not-in-family	White	Male	0	0
40						
3	Handlers-cleaners	Husband	Black	Male	0	0
40						
4	Prof-specialty	Wife	Black	Female	0	0
40						


	Country	Income_>50k
0	United-States	<=50K
1	United-States	<=50K
2	United-States	<=50K
3	United-States	<=50K
4	Cuba	<=50K

```
plt.figure()
df['Education'].value_counts().plot(kind='bar')
plt.title('Education Types')
plt.show()
```




```
sns.histplot(data=df, x='Age', bins=35, kde=True)
plt.show()
```

C:\Users\Muhammad Triski\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
 with pd.option\_context('mode.use\_inf\_as\_na', True):

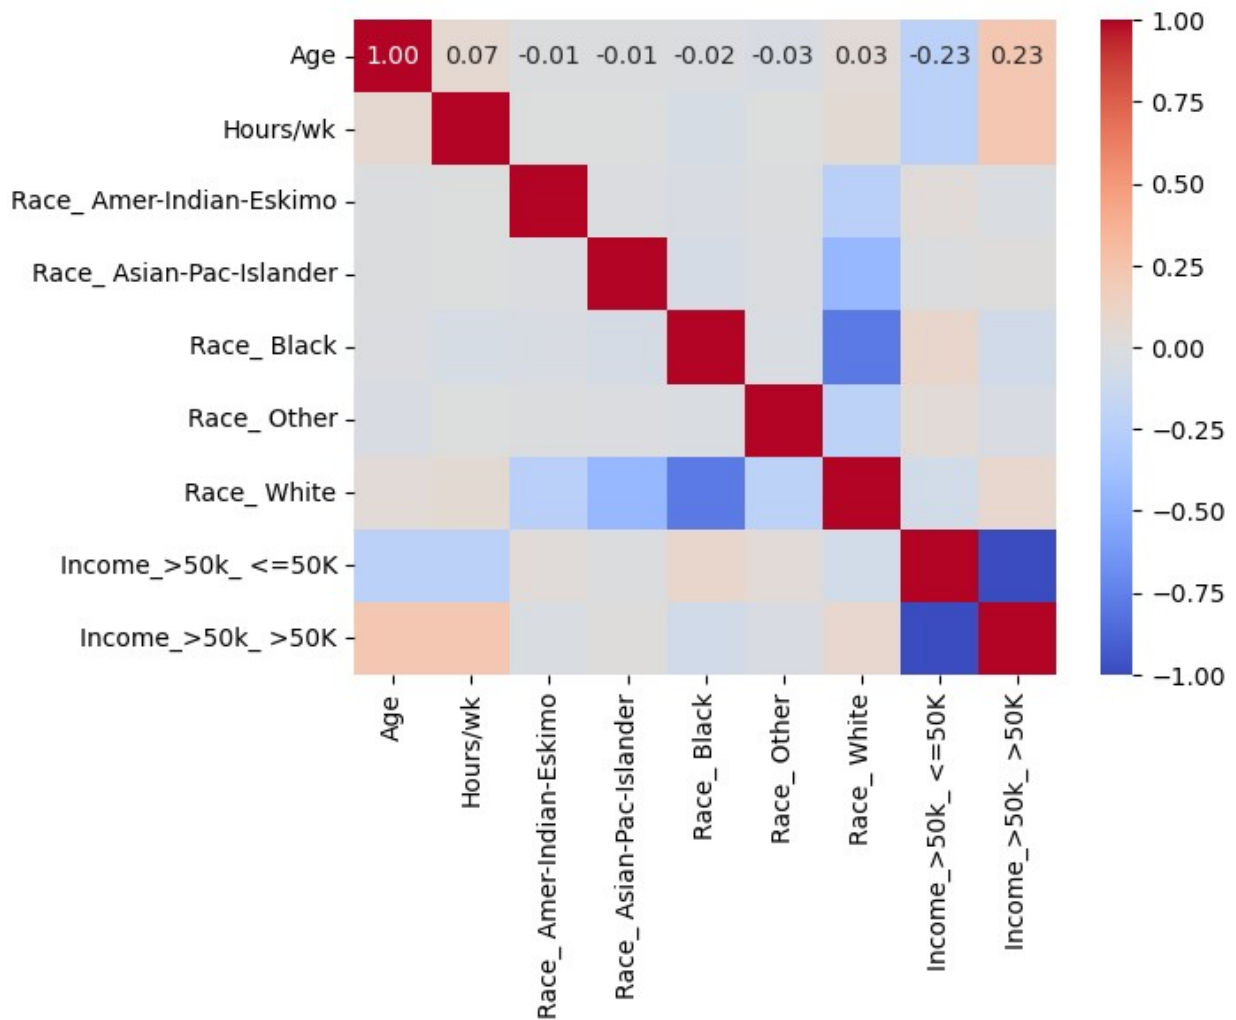


```
sns.histplot(data=df, x='idk', bins=35, kde=True)
plt.show()
```

```
C:\Users\Muhammad Triski\anaconda3\Lib\site-packages\seaborn\
_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated
and will be removed in a future version. Convert inf values to NaN
before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```



```
df_temp = pd.get_dummies(df[['Age', 'Race', 'Hours/wk',  
'Income_>50k']])  
  
corr_matrix = df_temp.corr()  
  
plt.figure()  
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm',  
cbar=True)  
plt.show()
```



```
df.replace('?', None, inplace=True)
df.isnull().sum()
```

```
Age      0
Sector   1836
idk      0
Education  0
idk2     0
Marital_Stat  0
Job      1843
Family   0
Race     0
Gender   0
idk3     0
idk4     0
Hours/wk  0
Country  583
```

```

Income_>50k      0
dtype: int64

df.dropna(inplace=True)
try:
    df = df.drop(columns=['idk2', 'idk3', 'idk3', 'idk4'])
except KeyError:
    pass

encoder = LabelEncoder()
df['Income_>50k'] = encoder.fit_transform(df['Income_>50k'])

df['Sector'] = df['Sector'].astype('category').cat.codes
df['Education'] = df['Education'].astype('category').cat.codes
df['Marital_Stat'] = df['Marital_Stat'].astype('category').cat.codes
df['Job'] = df['Job'].astype('category').cat.codes
df['Family'] = df['Family'].astype('category').cat.codes
df['Race'] = df['Race'].astype('category').cat.codes
df['Gender'] = df['Gender'].astype('category').cat.codes
df['Country'] = df['Country'].astype('category').cat.codes

scaler = StandardScaler()
attrs = df.columns[df.columns != 'Income_>50k']
df[attrs] = scaler.fit_transform(df[attrs])

df.head()

```

	Age	Sector	idk	Education	Marital_Stat	Job
Family \						
0	0.042796	2.936000	-1.062722	-0.349865	0.947847	-1.479055
1	0.880288	1.887682	-1.007871	-0.349865	-0.387275	-0.734545
2	-0.033340	-0.208955	0.244693	0.174763	-1.722396	-0.238206
3	1.108695	-0.208955	0.425240	-2.448375	-0.387275	-0.238206
4	-0.794697	-0.208955	1.406658	-0.349865	-0.387275	0.754473

	Race	Gender	Hours/wk	Country	Income_>50k
0	0.385048	0.692806	-0.077734	0.264924	0
1	0.385048	0.692806	-2.331531	0.264924	0
2	0.385048	0.692806	-0.077734	0.264924	0
3	-2.011035	0.692806	-0.077734	0.264924	0
4	-2.011035	-1.443405	-0.077734	-5.304034	0

*#Random forest*

```

X = df.drop('Income_>50k', axis=1)
y = df['Income_>50k']

```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2)
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_model = RandomForestClassifier(n_estimators=200)  
rf_model.fit(X_train, y_train)
```

```
RandomForestClassifier(n_estimators=200)
```

```
y_pred = rf_model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)  
precision = precision_score(y_test, y_pred)  
recall = recall_score(y_test, y_pred)  
f1 = f1_score(y_test, y_pred)
```

```
print(f'Accuracy: {accuracy:.4f}')
```

```
print(f'Precision: {precision:.4f}')
```


```
print(f'Recall: {recall:.4f}')
```

```
print(f'F1 Score: {f1:.4f}')
```

```
cm = confusion_matrix(y_test, y_pred)
```


```
disp = ConfusionMatrixDisplay(cm)  
disp.plot()  
plt.show()
```

```
Accuracy: 0.8241  
Precision: 0.6853  
Recall: 0.5793  
F1 Score: 0.6278
```



```
feature_importances = rf_model.feature_importances_  
features = X.columns  
  
importance_df = pd.DataFrame({'Feature': features, 'Importance':  
feature_importances})  
  
plt.figure()  
plt.barh(importance_df['Feature'], importance_df['Importance'])  
plt.title('Feature Importances in Random Forest')  
plt.show()
```





```
import torch
print(torch.__version__)

2.3.1

import torch
import torch.nn as nn

X_train = torch.tensor(X_train.values, dtype=torch.float32)
y_train = torch.tensor(y_train.values, dtype=torch.float32)
X_train = X_train.clone().detach()
y_train = y_train.clone().detach()

X_test = torch.tensor(X_test.values, dtype=torch.float32)
y_test = torch.tensor(y_test.values, dtype=torch.float32)
X_test = X_train.clone().detach()
y_test = y_train.clone().detach()

X_train.shape

torch.Size([24129, 11])

class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
```

```

        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(11, 118),
            nn.ReLU(),
            nn.Linear(118, 118),
            nn.ReLU(),
            nn.Linear(118, 1)
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        logits = torch.sigmoid(logits)
        return logits

model = Model()
print(model)

Model(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=11, out_features=118, bias=True)
    (1): ReLU()
    (2): Linear(in_features=118, out_features=118, bias=True)
    (3): ReLU()
    (4): Linear(in_features=118, out_features=1, bias=True)
  )
)

criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

epochs = 200
for epoch in range(epochs):
    model.train()
    optimizer.zero_grad()
    outputs = model(X_train).squeeze()
    loss = criterion(outputs.squeeze(), y_train)
    loss.backward()
    optimizer.step()

    if (epoch + 1) % 10 == 0:
        print(f'Epoch [{epoch + 1}/{epochs}], Loss: {loss.item():.4f}')

Epoch [10/200], Loss: 0.5498
Epoch [20/200], Loss: 0.5055
Epoch [30/200], Loss: 0.4738
Epoch [40/200], Loss: 0.4590
Epoch [50/200], Loss: 0.4484

```

```
Epoch [60/200], Loss: 0.4397
Epoch [70/200], Loss: 0.4319
Epoch [80/200], Loss: 0.4247
Epoch [90/200], Loss: 0.4181
Epoch [100/200], Loss: 0.4124
Epoch [110/200], Loss: 0.4079
Epoch [120/200], Loss: 0.4046
Epoch [130/200], Loss: 0.4021
Epoch [140/200], Loss: 0.4002
Epoch [150/200], Loss: 0.3985
Epoch [160/200], Loss: 0.3971
Epoch [170/200], Loss: 0.3957
Epoch [180/200], Loss: 0.3943
Epoch [190/200], Loss: 0.3930
Epoch [200/200], Loss: 0.3918
```

```
model.eval()
y_true = []
y_pred = []

with torch.no_grad():
    outputs = model(X_test)
    predicted = outputs.round()

    y_true = y_test.cpu().numpy()
    y_pred = predicted.cpu().numpy()

y_true = np.array(y_true).flatten()
y_pred = np.array(y_pred).flatten()

accuracy = accuracy_score(y_true, y_pred)
precision = precision_score(y_true, y_pred)
recall = recall_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred)

print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1 Score: {f1:.4f}')

cm = confusion_matrix(y_true, y_pred)

disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.show()

Accuracy: 0.8046
Precision: 0.6152
Recall: 0.5591
F1 Score: 0.5858
```

